

TP4 :HBASE

```
PS C:\Users\DELL\Downloads\HBase> docker start hadoop-master hadoop-slave1 hadoop-slave2  
hadoop-master  
hadoop-slave1  
hadoop-slave2  
PS C:\Users\DELL\Downloads\HBase> docker exec -it hadoop-master bash  
root@hadoop-master:~# ./start-hadoop.sh  
  
Starting namenodes on [hadoop-master]  
hadoop-master: Warning: Permanently added 'hadoop-master,172.21.0.2' (ECDSA) to the list of known hosts.  
  
root@hadoop-master:~# /usr/local/hbase/bin/start-hbase.sh  
hadoop-master: Warning: Permanently added 'hadoop-master,172.21.0.2' (ECDSA) to the list of known hosts.  
hadoop-master: running zookeeper, logging to /usr/local/hbase/bin/../logs/hbase-root-zookeeper-hadoop-master.out  
root@hadoop-master:~# jps  
1536 HQuorumPeer  
1786 HRegionServer  
251 NameNode  
1611 HMaster  
539 SecondaryNameNode  
844 ResourceManager  
2220 Jps
```

```
root@hadoop-master:~# hbase shell  
SLF4J: Class path contains multiple SLF4J bindings.
```

- Nous allons créer une base de données qui contient les données suivantes:

Row Key	customer		sales	
	ROW_ID	name	city	product
101	John White	Los Angeles, CA	Chairs	\$400.00
102	Jane Brown	Atlanta, GA	Lamps	\$200.00
103	Bill Green	Pittsburgh, PA	Desk	\$500.00
104	Jack Black	St. Louis, MO	Bed	\$1,600.00

```
'ASSI' hbase(main):001:0> create 'sales_ledger', 'customer', 'sales'  
0 row(s) in 1.9120 seconds  
  
=> Hbase::Table - sales_ledger  
hbase(main):002:0> █
```

```

hbase(main):001:0> create 'sales_ledger','customer','sales'
0 row(s) in 1.9120 seconds

=> Hbase::Table - sales_ledger
hbase(main):002:0> list
TABLE
sales_ledger
1 row(s) in 0.0310 seconds

=> ["sales_ledger"]
hbase(main):003:0>

```

Chargement :

```

PS C:\Users\DELL\Downloads\HBase> docker exec hadoop-master bash -c "hbase org.apache.hadoop.hbase.ma
'preduce.ImportTsv -libjars /usr/local/hbase/lib/commons-lang-2.6.jar -Dimportttsv.separator=',' -Dimpo
rttsv.columns=HBASE_ROW_KEY,cf:date,cf:time,cf:town,cf:product,cf:price,cf:payment products input/pur
chases2.txt"

```

Insérer les différentes lignes:

```

hbase(main):012:0> put 'sales_ledger','103','customer:name','Bill Green'
'sburgh, PA'
put 'sales_ledger','103','sales:product',0 row(s) in 0.0150 seconds

't','Desk'
put 'sales_ledger','103','sales:amount',hbase(main):013:0> put 'sales_ledger','103','customer:city','Pitts
0 row(s) in 0.0080 seconds

hbase(main):014:0> put 'sales_ledger','103','sales:product','Desk'
0 row(s) in 0.0070 seconds

```

Visualiser le résultat de l'insertion, en tapant:

```

hbase(main):020:0> scan 'sales_ledger'
ROW                                COLUMN+CELL
 101                               column=customer:city, timestamp=1763028623521, value=Los Angeles, CA
 101                               column=customer:name, timestamp=1763028623458, value=John White
 101                               column=sales:amount, timestamp=1763028625894, value=$400.00
 101                               column=sales:product, timestamp=1763028623558, value=Chairs
 102                               column=customer:city, timestamp=1763028660071, value=Atlanta, GA
 102                               column=customer:name, timestamp=1763028659969, value=Jane Brown
 102                               column=sales:amount, timestamp=1763028661478, value=$200.00
 102                               column=sales:product, timestamp=1763028660172, value=Lamps
 103                               column=customer:city, timestamp=1763028671158, value=Pittsburgh, PA

```

Afficher les valeurs de la colonne product de la ligne 102

```

hbase(main):021:0> get 'sales_ledger','102',{COLUMN => 'sales:product'}
COLUMN                                CELL
  sales:product                         timestamp=1763028660172, value=Lamps
1 row(s) in 0.0390 seconds

```

Explorer les filtres dans HBase :

Filter – RowFilter

```
hbase(main):006:0> scan 'sales_ledger', {FILTER => "RowFilter(>, 'binary:102')"}  
ROW  
      COLUMN+CELL  
103          column=customer:city, timestamp=1763028671158, value=Pittsburgh, PA  
103          column=customer:name, timestamp=1763028671118, value=Bill Green  
103          column=sales:amount, timestamp=1763028673600, value=$500.00  
103          column=sales:product, timestamp=1763028671188, value=Desk  
104          column=customer:city, timestamp=1763028694679, value=St. Louis, MO  
104          column=customer:name, timestamp=1763028694615, value=Jack Black  
104          column=sales:amount, timestamp=1763028705259, value=$1,000.00  
104          column=sales:product, timestamp=1763028694713, value=Bed  
2 row(s) in 0.0670 seconds
```

```
hbase(main):007:0> scan 'sales_ledger', {FILTER => "FamilyFilter(=, 'binary:customer')}'  
ROW  
      COLUMN+CELL  
101          column=customer:city, timestamp=1763028623521, value=Los Angeles, CA  
101          column=customer:name, timestamp=1763028623458, value=John White  
102          column=customer:city, timestamp=1763028660071, value=Atlanta, GA  
102          column=customer:name, timestamp=1763028659969, value=Jane Brown  
103          column=customer:city, timestamp=1763028671158, value=Pittsburgh, PA  
103          column=customer:name, timestamp=1763028671118, value=Bill Green  
104          column=customer:city, timestamp=1763028694679, value=St. Louis, MO  
104          column=customer:name, timestamp=1763028694615, value=Jack Black  
4 row(s) in 0.0310 seconds
```

Filter – FamilyFilter

```
hbase(main):008:0> scan 'sales_ledger', {FILTER => "FamilyFilter(=, 'binary:customer')}  
ROW  
      COLUMN+CELL  
101          column=customer:city, timestamp=1763028623521, value=Los Angeles, CA  
101          column=customer:name, timestamp=1763028623458, value=John White  
102          column=customer:city, timestamp=1763028660071, value=Atlanta, GA  
102          column=customer:name, timestamp=1763028659969, value=Jane Brown  
103          column=customer:city, timestamp=1763028671158, value=Pittsburgh, PA  
103          column=customer:name, timestamp=1763028671118, value=Bill Green  
104          column=customer:city, timestamp=1763028694679, value=St. Louis, MO  
104          column=customer:name, timestamp=1763028694615, value=Jack Black  
4 row(s) in 0.0320 seconds
```

Filter – QualifierFilter

```
hbase(main):009:0> scan 'sales_ledger', {FILTER => "QualifierFilter(=, 'binary:amount')}  
ROW  
      COLUMN+CELL  
101          column=sales:amount, timestamp=1763028625894, value=$400.00  
102          column=sales:amount, timestamp=1763028661478, value=$200.00  
103          column=sales:amount, timestamp=1763028673600, value=$500.00  
104          column=sales:amount, timestamp=1763028705259, value=$1,600.00  
4 row(s) in 0.0280 seconds
```

Filter – SingleColumnValueFilter

```
hbase(main):010:0> scan 'sales_ledger', {FILTER => "SingleColumnValueFilter('sales', 'amount', =, 'binary:  
hbase(main):011:1" $450.00')"}  
ROW  
      COLUMN+CELL  
0 row(s) in 0.0290 seconds
```

Filter – PrefixFilter

```
hbase(main):001:0> scan 'sales_ledger', {FILTER => "PrefixFilter('10')"}  
ROW  
      COLUMN+CELL  
101          column=customer:city, timestamp=1763028623521, value=Los Angeles, CA  
101          column=customer:name, timestamp=1763028623458, value=John White  
101          column=sales:amount, timestamp=1763028625894, value=$400.00  
101          column=sales:product, timestamp=1763028623558, value=Chairs  
102          column=customer:city, timestamp=1763028660071, value=Atlanta, GA  
102          column=customer:name, timestamp=1763028659969, value=Jane Brown  
102          column=sales:amount, timestamp=1763028661478, value=$200.00  
102          column=sales:product, timestamp=1763028660172, value=Lamps
```

Filter – PageFilter

```
hbase(main):002:0> scan 'sales_ledger', {FILTER => "PageFilter(2)"}
ROW                                COLUMN+CELL
101                               column=customer:city, timestamp=1763028623521, value=Los Angeles, CA
101                               column=customer:name, timestamp=1763028623458, value=John White
101                               column=sales:amount, timestamp=1763028625894, value=$400.00
101                               column=sales:product, timestamp=1763028623558, value=Chairs
102                               column=customer:city, timestamp=1763028660071, value=Atlanta, GA
102                               column=customer:name, timestamp=1763028659969, value=Jane Brown
102                               column=sales:amount, timestamp=1763028661478, value=$200.00
102                               column=sales:product, timestamp=1763028660172, value=Lamps
2 row(s) in 0.0440 seconds
```

Chargement donne :

```
PS C:\Users\DELL\Downloads\HBase> docker exec hadoop-master bash -c 'export HADOOP_CLASSPATH=$(hbase classpath) && hbase org.apache.hadoop.hbase.mapreduce.ImportTsv -Dimporttsv.separator="," -Dimporttsv.columns=HBASE_ROW_KEY,cf:date,cf:time,cf:town,cf:product,cf:price,cf:payment products /user/root/input/purchases2.txt'
products /user/root/input/purchases2.txt';d79dace4-461b-421c-9a6a-7c72ef73000aSLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/local/hbase/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/local/hadoop/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/
PS C:\Users\DELL\Downloads\HBase> docker cp "C:\Users\DELL\Downloads\HBase\hbase-code\scan.hbase" hadoop-master:/tmp/scan.hbase; docker exec hadoop-master bash -c "hbase shell /tmp/scan.hbase 2>&1" | Select-String -Pattern "ROW|column=|row\(\s\)"
461b-421c-9a6a-7c72ef73000aSuccessfully copied 2.05kB to hadoop-master:/tmp/scan.hbase

ROW COLUMN+CELL
1 column=cf:date, timestamp=1763033080722, value=2012-01-01
1 column=cf:payment, timestamp=1763033080722, value=Amex
1 column=cf:price, timestamp=1763033080722, value=214.05
1 column=cf:product, timestamp=1763033080722, value=Men's clothing
1 column=cf:time, timestamp=1763033080722, value=09:00
```